

Myntra StyleGenie — What I Built & How

A GenAI-native styling layer reimagining Myntra: conversational discovery, visual search, generative outfits, AI fit prediction and a personalised trend digest — layered on top of the existing app experience.

🔗 **Live prototype:** <https://style-genie-myntra.vercel.app/>

React + Vite prototype

5 GenAI features

Build ✓ Lint ✓ 23/23 tests ✓

8-page strategy deck


This document explains the full submission end-to-end so it can be presented confidently in an interview. It covers the prototype, each feature, the architecture, how the AI is simulated vs. how it would run in production, the testing performed, the business case, and a Q&A bank mapped to the evaluation rubric.

1 · The submission at a glance


Deliverable	What it is	File / location
Prototype (live)	A working Myntra-styled web app with all five StyleGenie GenAI features wired in and interactive. Live: https://style-genie-myntra.vercel.app/	stylegenie/ (run <code>npm run dev</code> ; deployed on Vercel)
Strategy Deck	8-page deck (cover + 7): Problem · Why GenAI · Approach & Solution · Implementation Flowchart · Live Prototype · Metrics · Risks.	StyleGenie_Strategy_Deck.pdf
Build report	This document — what was built, how, testing, and interview prep.	StyleGenie_Build_Report.pdf

2 · What the prototype is


A single-page React application that **looks like Myntra** (logo, pink accent, category nav, product grid, product detail page) with a GenAI layer woven in — exactly as the brief asks ("the UI should look like the existing product itself"). It is fully self-contained: every product image is a generated, branded SVG, and every "AI" response is produced deterministically on the client, so the demo works offline, never shows a broken image, and needs no API keys.




Home — Myntra-style storefront with an "Ask StyleGenie" concierge and AI-styling entry points.




StyleGenie Chat — natural-language request parsed into intent tags, grounded product cards streamed back.



Product page — FitGenius (AI size) and Outfit Studio (AI stylist) embedded inline.



TrendPulse — AI-generated daily trend narratives, each with curated catalogue products.



Visual Match — upload an image / pick a vibe → CLIP-style visual search returns ranked matches.

3 · The five GenAI features (what / how / why GenAI / impact)

① StyleGenie Chat — AI style concierge

What: a floating concierge; users type what they need in plain language. **How:** message → intent classifier extracts occasion, budget, gender & style → retrieval scores the catalogue → curated outfit cards stream back with the parsed intent shown as tags. **Why GenAI:** only an LLM can parse "*something Alia wore but more affordable, under 4K*" — four constraints in one sentence. **Impact:** collapses 40-product browsing into ~5 curated picks; target 12% chat-to-cart.

② Visual Match — snap & discover

What: upload any photo / screenshot / Pinterest save and find visually similar Myntra products. **How:** image → CLIP embedding → nearest-neighbour against pre-computed catalogue vectors → ranked matches with % scores. **Why GenAI:** CLIP understands aesthetics ("boho", "quiet luxury") that colour-histogram CV cannot. **Impact:** ~35% search-to-PDP CTR vs keyword search.

③ Outfit Studio — AI outfit composer

What: pick any product, get 3–5 complete, colour-coordinated looks built from real catalogue items, each tagged to an occasion, with a one-tap "add complete look to bag". **How:** anchor product → role-aware composition (top/bottom/footwear/accessory) scored on shared style + occasion fit. **Why GenAI:** composing a coherent outfit needs creative judgement, not a similarity lookup. **Impact:** +1.8 items/order → ~22% higher AOV.

④ FitGenius — smart fit predictor

What: enter a quick body profile, get your true size with a confidence score and the *evidence* behind it. **How:** body profile × NLP mining of review text ("runs slim", "true to size") → size + confidence + sources. **Why GenAI:** static size charts are lookup tables; an LLM turns thousands of unstructured reviews into a fit signal. **Impact:** return rate 32% → <18% (~₹525 Cr/yr saved) — the single biggest lever.

⑤ TrendPulse — AI trend digest

What: a personalised daily lookbook. **How:** social signals → LLM writes a trend narrative → products curated against that trend *and* the user's style graph. **Why GenAI:** generating fresh narrative content per user per day is inherently generative. **Impact:** +40% DAU from a daily check-in habit.

4 · Architecture & tech stack

Layer	In the prototype	In production
Presentation	React 19 + Vite, hand-written CSS in Myntra's design language; floating concierge + inline AI cards.	Same UI shell on the live Myntra app.
AI orchestration	<code>src/ai/engine.js</code> — intent parsing, retrieval/scoring, outfit composition, fit prediction, visual match, trend generation.	LLM router → intent classifier → feature dispatcher → response generator.
Models	Deterministic stand-ins for each model call.	GPT-4o / Gemini (chat), CLIP (visual), fit model over reviews, diffusion (future try-on).
Data	<code>src/data/products.js</code> — 24-SKU mock catalogue; branded SVG images.	Product catalogue embeddings, user style graph, review corpus, trend signals (vector DB + RAG).

How the AI is simulated (be transparent about this in the interview): the prototype's "AI" runs entirely client-side and deterministically — intent is parsed with keyword/occasion/budget extraction, retrieval is a transparent scoring function, and outfits are composed by role. This makes the demo fast, free and reliable. The exact same call sites (`parseIntent` , `searchCatalog` , `composeOutfits` , `predictFit` , `visualMatch` , `trendDigest`) are where a real backend would

orchestrate LLM + CLIP + a RAG pipeline. **The key production principle is RAG grounding:** the LLM provides intelligence, the live catalogue provides facts — so it can never recommend a product that doesn't exist.

5 · Testing performed (error-free, verified)

- ✓ **Production build** — `npm run build` compiles 25 modules with zero errors/warnings.
- ✓ **Lint** — `npm run lint` (ESLint, React 19 rules) passes clean; fixed an effect-ordering issue, a `useState-in-effect` warning, an unused variable, and a fast-refresh export rule.
- ✓ **AI engine unit tests** — 23/23 assertions pass: intent parsing ("college fest under 2K" → occasion=college, budget=2000), budget-respecting retrieval, multi-item outfit composition with correct totals, dresses correctly skip a bottom, fit sizes/confidence in valid ranges, descending visual-match scores, 5 trends with products, and all 24 product images valid data-URLs.
- ✓ **Render smoke test** — headless Chrome renders the app with 0 runtime errors; verified Home, Product Detail, TrendPulse, Chat (with a live query returning grounded cards) and Visual Match all mount and display correctly.
- ✓ **Bug found & fixed** — the "H&M" brand put a raw `&` into an SVG (invalid XML), breaking one product image; fixed with XML-escaping + standard data-URI encoding.

6 · Business case

Feature	Primary metric	Target
StyleGenie Chat	Chat-to-cart conversion	12% in 3 months
Visual Match	Search-to-PDP CTR	35%
Outfit Studio	Items per order	+1.8 (≈ +22% AOV)
FitGenius	Return rate	<18% (from 32%)
TrendPulse	DAU	+40%

Projected annual impact: ₹800–1200 Cr revenue uplift + ₹400–525 Cr cost savings + ₹50–80 Cr premium subscriptions = **₹1200–1700 Cr net**. Validated through a 50/50 A/B holdout run over 4–6 weeks (conversion, return rate, AOV, items/order, 7-day retention, feature-specific funnels).

7 · Pitfalls & how we handle them

Risk	Mitigation
Hallucinations	RAG grounding in the live catalogue; LLM never invents SKUs; confidence scoring filters low-certainty output.
Privacy	On-device photo processing, explicit opt-in, transparent data policy, no third-party sharing.
Latency	Streaming responses, pre-computed embedding caches, edge-deployed CLIP; target <2s.
User trust	Confidence scores + review evidence + manual override at every step.
Scalability / cost	Tiered models (cheap→premium), aggressive caching, batch for non-realtime (TrendPulse).
Adoption	Soft inline integration, contextual nudges on struggle signals, free tier + premium upsell.

8 · Interview Q&A bank

Q. Why Myntra?

Fashion e-commerce has the most painful, GenAI-solvable problems (30–40% returns, ~70% cart abandonment, choice overload), is inherently visual & contextual (perfect for multimodal AI), and the impact is large and measurable. Unlike food delivery where the core flow is already efficient, fashion discovery is fundamentally broken.

Q. How is this different from Myntra's current recommendations?

Today's recos are predictive ML — collaborative filtering ("users who bought X bought Y"), backward-looking and generic. StyleGenie understands *intent*: "beach vacation, casual but Instagram-worthy, under 5K" parses four constraints at once. Collaborative filtering can't.

Q. How do you prevent hallucinations?

Every response is RAG-grounded: the LLM retrieves from the live catalogue and never generates product facts from training data. Intelligence from the LLM, facts from the catalogue, confidence scoring on top.

Q. Why not just use ChatGPT directly?

A raw LLM doesn't know Myntra's catalogue and would recommend items that don't exist. The innovation is the integration layer — grounding the LLM in real-time product data via embeddings + RAG. It's an AI system, not a chatbot.

Q. What about cost at scale?

Tiered models (small/cheap for simple queries, premium for complex styling), aggressive caching (the first "Diwali outfit under 5K" is cached and personalised cheaply), and batch pre-computation for TrendPulse.

Q. How did you build the prototype?

A React + Vite app replicating Myntra's design language. Five AI features as interactive overlays; chat streams simulated LLM responses. Every AI call site maps 1:1 to where a production backend would orchestrate the real models, so it's a faithful blueprint, not a mockup.

Q. What metrics would you A/B test?

50/50 holdout for 4–6 weeks. Primary: conversion, return rate, AOV, items/order, 7-day retention. Feature-specific: chat-to-cart, visual-search CTR.

9 · Mapping to the evaluation rubric

Evaluation parameter	How this submission addresses it
Problem clarity & depth	Four quantified pain points tied directly to revenue (slide 1).
Suitability of GenAI usage	Each feature justified vs. the traditional ML alternative; strictly generative/multimodal, no predictive ML (slide 2).
Quality of solutioning	Five features spanning intent → discovery → confidence → checkout, with an orchestration layer.
Uniqueness & originality	Occasion-first framing + generative outfit composition + evidence-backed fit, not just a search box.
Business impact potential	Per-feature metrics and a ₹1200–1700 Cr modelled annual impact.
Deck flow & storytelling	Cover → problem → why GenAI → approach & solution → implementation flowchart → live prototype → metrics → risks.
Exhaustiveness of design	Architecture, data layer, pitfalls + mitigations, and a production path all specified.
Prototype functionality	All five features interactive; build + lint + 23 tests pass.

10 · How to run & deploy

Run locally: `cd stylegenie` → `npm install` → `npm run dev` → open the printed localhost URL.

Production build: `npm run build` (outputs static files to `stylegenie/dist/`).

Published (Part 2 live URL): <https://style-genie-myntra.vercel.app/> — deployed on Vercel from the GitHub repo, zero config, no env vars. To redeploy or host elsewhere, the static build works in one step on any accepted platform:

- **Vercel:** `cd stylegenie && npx vercel --prod` (or import the folder at vercel.com).
- **Netlify:** drag the `stylegenie/dist` folder onto app.netlify.com/drop, or `npx netlify deploy --prod --dir=dist`.
- **Bolt / Lovable / Replit:** import the `stylegenie` folder; framework = Vite, build = `npm run build`, output = `dist`.

Deep links for demoing specific states: `#trends`, `#p/12` (a product page), `#genie` (chat), `#genie/festive%20look%20under%203k` (auto-runs a chat query), `#visual` (Visual Match) — append these to the live URL to demo each feature instantly.