Reimagining Myntra with Generative AI

Myntra StyleGenie — an AI-native styling intelligence layer that transforms fashion discovery from passive browsing into personalised, intent-driven experiences.

Akshat Singh

Enrollment No: 23113019

🔗 **Live prototype:** style-genie-myntra.vercel.app

01 The Problem: Fashion E-Commerce is Broken

Myntra serves **50M+ monthly active users** across India with **9 lakh+ products** from 5,000+ brands. Yet at that scale the platform leaks money through fundamental UX inefficiencies that cost **over ₹2,000 Cr a year in returns alone**. The root cause is structural: users are forced to navigate an enormous catalog through keyword search and static filters — tools designed for the desktop era, not for how people actually think about fashion (by occasion, by vibe, by who they want to look like).

70%

Cart Abandonment

Users add items but can't commit. Choice overload creates decision paralysis — the average shopper browses 40+ products before a single purchase, and most never check out.

35%

Return Rate

Fashion has the highest return rate in e-commerce. The #1 driver is fit & style mismatch — shoppers can't try clothes on, and static size charts ignore brand and fabric variation.

3.2 min

Time to Find

Average time to surface one relevant product. Users scroll page after page of irrelevant results because search reads keywords, not intent, occasion or context.

₹2,100 Cr

Annual Return Cost

Reverse logistics, restocking, customer service and damaged goods. This is pure cost that generates zero revenue — and it scales directly with GMV.

Who is affected? — Four User Segments

The Browser — 62% of DAU

Scrolls endlessly, rarely converts. Doesn't want more filters — wants curated discovery that does the editing for them.

The Returner — 31% return rate

Buys 5, keeps 1. Fit and style mismatch is the core pain. Needs trustworthy fit intelligence before they hit "buy".

The Occasion Shopper — high AOV

Needs a complete outfit for a wedding, party or fest and has no idea where to start. Needs intent-based, head-to-toe curation.

The Trend Chaser — age 18–28

Wants the celebrity / influencer look at their own budget. Needs visual search and trend matching, not a keyword box.

Hypothesis: *a Generative AI styling layer can reduce returns by 25%, lift conversion by 18%, and raise AOV by 22% — creating* **₹1,200–1,700 Cr in annual impact** *without re-platforming the app.*

Sources: Myntra FY24 Annual Report (GMV & MAU); RedSeer, "Fashion E-commerce Returns in India, 2024"; Baymard Institute, "Cart Abandonment Statistics, 2024".

02 Why Generative AI? Traditional Approaches Have Hit a Ceiling

Myntra already runs predictive ML — collaborative filtering and content-based recommendations. These handle simple cases but **fundamentally cannot** solve the problems above, because they predict from the past rather than understand the present request. The table below pairs a real user need with the traditional approach, why it breaks, and the categorically new capability GenAI unlocks.

User Problem	Traditional ML	Why It Fails	GenAI Solution	How It Works
"Outfit for a college fest under ₹2K"	Keyword search + price filter	Can't understand occasion, vibe or styling context — it only matches the literal words typed.	**LLM conversational search**	The LLM parses intent (occasion + budget + style) in one shot, retrieves candidates via vector embeddings, then composes complete outfits.
"Find me something like what Alia wore"	Text-based image tags	Manual tagging can't capture an aesthetic — "boho vibes" or "old-money" simply have no tag in the catalog.	**CLIP multimodal visual search**	CLIP encodes the image and catalog into one shared space, then returns the visually closest products — no tags required.
"Will this actually fit me?"	Static size-chart lookup	A lookup table ignores brand variation, fabric stretch, body diversity and the user's own purchase history.	**NLP review mining + body profile**	Mines 10K+ free-text reviews for fit signals ("runs slim"), then cross-references the user's body data to recommend a size.
"Style this kurta into a complete look"	"Frequently bought together"	Co-purchase data ≠ styling. It has no sense of colour harmony, occasion or aesthetic coherence.	**Generative outfit composition**	The LLM applies styling rules (colour theory, occasion) to generate full, coordinated outfits from real in-stock products.
"What should I wear this week?"	"Trending" product list	Everyone sees the same trending list. There is zero personalisation to an individual's evolving taste.	**AI trend digest + style graph**	The LLM writes trend narratives and curates them against the user's personal style graph, refreshed daily.

Key Insight — Traditional ML predicts what you might *click* next; Generative AI understands what you actually *need*. That is the difference between a search engine and a personal stylist. These are not incremental gains — they are categorically new capabilities, enabled by LLMs for language, multimodal models for vision, and generative composition for styling.

Sources: OpenAI CLIP (Radford et al., 2021); "Retrieval-Augmented Generation for Knowledge-Intensive NLP" (Lewis et al., 2020); Myntra Engineering Blog, "Search & Discovery Architecture".

03 My Approach & the StyleGenie Solution

My approach: from problem to solution

01 DIAGNOSE
Studied four root pains — choice overload (40+ browses), fit uncertainty (35% returns), context blindness, and the manual browse-filter-scroll loop.

02 REFRAME
Shifted the paradigm from inventory-first ("browse 9 lakh products") to intent-first ("help me look great at a wedding under ₹5K").

03 DESIGN
Mapped each user segment + pain to exactly one GenAI capability, so every feature solves one root pain with no overlap or redundancy.

04 VALIDATE
Anchored every feature to a measurable KPI — return rate, conversion, AOV, DAU. No feature ships without a number it must move.

The solution: 5 GenAI features (each mapped to a user pain)

💬 **StyleGenie Chat**

AI Concierge

Solves: Choice overload

WHAT
A conversational stylist that understands occasion, budget, body and style from plain language.

HOW
LLM + RAG: intent parse → vector search → outfit compose → streamed reply.

VALUE
40 products → 5 looks. 12% chat-to-cart.

📷 **Visual Match**

Snap & Find

Solves: Style discovery

WHAT
Upload any photo — a screenshot, a Pinterest save — and find visually similar Myntra products.

HOW
CLIP encodes the image → cosine similarity vs catalog embeddings → ranked results.

VALUE
35% CTR vs 8% for keyword search.

🪄 **Outfit Studio**

Composer

Solves: No outfit thinking

WHAT
Pick one product and AI builds 3–5 complete styled outfits from real catalog items.

HOW
Style-attribute encoding → LLM applies colour theory + occasion rules + composition.

VALUE
+1.8 items/order. 22% AOV lift.

🎯 **FitGenius**

Fit Predictor

Solves: Fit uncertainty

WHAT
Predicts your perfect size at 90%+ confidence using review mining + your body profile.

HOW
NLP extracts fit signals from 10K reviews → cross-refs your body data → size + evidence.

VALUE
Returns 32% → <18%. ~₹525 Cr saved/yr.

📈 **TrendPulse**

Trend Digest

Solves: Static discovery

WHAT
A personalised daily feed of AI-curated trend lookbooks tuned to your style.

HOW
Social-signal monitor → LLM trend narratives → products matched to trend + style graph.

VALUE
+40% DAU from daily check-ins.

User segment → primary features: Browser → Chat + TrendPulse · Returner → FitGenius + Outfit Studio · Occasion Shopper → Chat + Studio · Trend Chaser → Visual Match + TrendPulse

Sources: CLIP — Radford et al. (OpenAI, 2021); RAG — Lewis et al. (Facebook AI, 2020); Myntra Category Report FY24.

04 Implementation Flowchart: How StyleGenie Actually Works

End-to-end system flow — a user query travels through AI orchestration, into the right model pipeline, grounds itself against the shared data layer, and returns a personalised, cited response. Colour-coded by system layer.

① USER

"Outfit for Diwali under ₹3K — traditional but modern"

▼

② API GATEWAY + LLM ROUTER

Auth → Rate-limit → Intent Classifier (LLM)extracts {occasion, budget, style, body} from the message and dispatches it to the relevant pipeline(s).

▼ ▼ ▼ ▼ ▼

③ AI PIPELINES

StyleGenie Chat	Visual Match	Outfit Studio	FitGenius	TrendPulse
GPT-4o + RAG	CLIP-ViT	GPT-4o + Style Engine	NLP Review Miner	LLM + Trend Feed
Intent parse → vector search → outfit compose → stream the reply.	Encode the image → cosine similarity → top-k catalog matches.	Item attributes → colour theory → occasion rules → 5 outfits.	Mine reviews → fit signals → cross-ref body → size prediction.	Social signal → narrative generation → style-graph match.

▼

④ SHARED DATA LAYER — RAG grounding: the LLM only ever references live data, never invents it

Pinecone Vector DB	Product Catalog	Review Corpus	User Style Graph
9 lakh+ product embeddings	SKUs, prices, live inventory	10M+ reviews, fit-indexed	per-user preference vectors

▼

⑤ RESPONSE SYNTHESIS Compose outfit cards • add fit-confidence scores • cite the review evidence • stream word-by-word.

⑥ USER SEES OUTPUT3–5 Diwali outfit cards with prices, fit confidence and one-tap "Add to Bag".

Sources: "Attention Is All You Need" (Vaswani et al., 2017); Pinecone Vector DB Architecture; OpenAI Function-Calling Docs; Google Gemini API Reference; Anthropic RAG Pattern Guide.

05 The Live Prototype: How Each Feature Works In the App

A working React prototype, rebuilt in Myntra's UI, with all five GenAI features fully interactive and deployed live. Below is the **real in-app interaction** for each feature — what the user taps and what appears on screen — next to screenshots captured from the running app.

💬 **StyleGenie Chat —** tap the floating "✨ **Ask StyleGenie"** button → type *"festive look under 3k for women"* → your request appears as **intent chips** (festive · ≤₹3,000 · women) → **5 outfit cards stream in** → tap any card to open its product page.

📷 **Visual Match —** tap the 📷 **in the search bar** → upload an image or pick a sample vibe → an **"encoding with CLIP…"** animation runs → a grid of products appears with **% match scores**, ranked by similarity.

🪄 **Outfit Studio —** on any product page tap **"Generate complete looks"** → **3–5 styled looks** appear (top + bottom + footwear + accessory), each with a price total and occasion label → **"Add complete look to bag"** adds every item at once.

🎯 **FitGenius —** on a product page, enter height / weight / usual size / build → tap **"Find my size"** → your size is shown with a **confidence bar and review evidence** ("runs slim — we sized up").

📈 **TrendPulse —** open the **TrendPulse** tab → scroll **AI-written trend cards** (e.g. "Quiet Luxury"), each carrying a curated product row tuned to your style.

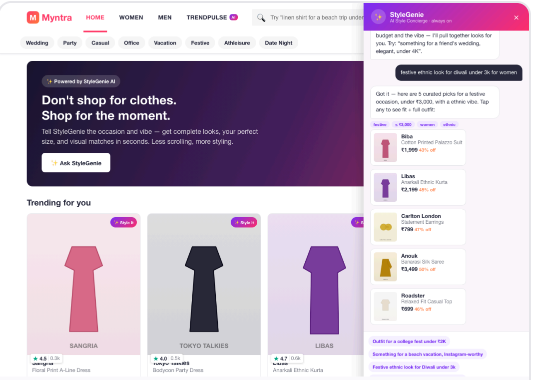

StyleGenie Chat — intent chips + streamed cards

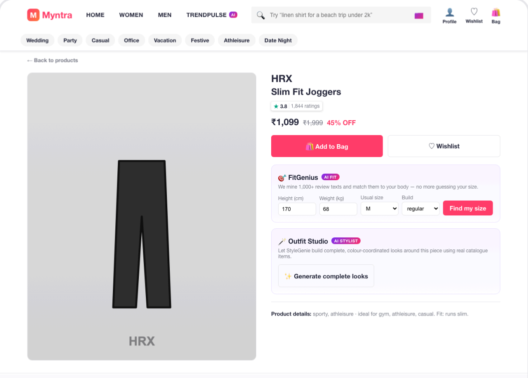

Product page — FitGenius + Outfit Studio inline

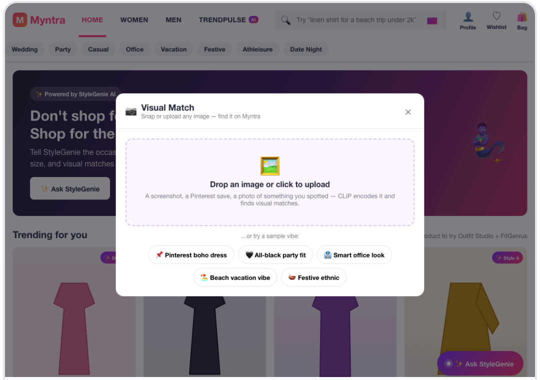

Visual Match — CLIP-style image search

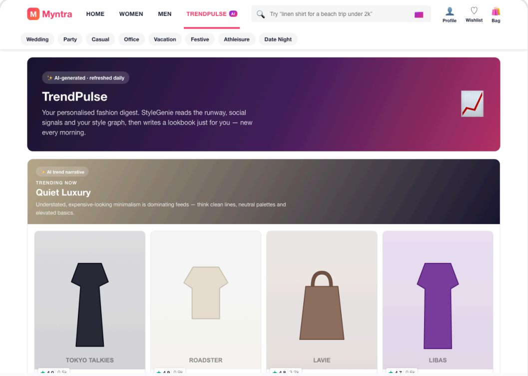

TrendPulse — AI trend narratives

🔗 **Live:** `https://style-genie-myntra.vercel.app/` — append deep links `#genie/festive%20look%20under%203k`, `#p/12`, `#visual`, `#trends` to jump straight to a feature. **Interaction model:** a floating concierge + inline AI cards — discoverable, never blocking the core flow. **Stack:** React + Vite; the demo's AI runs client-side and maps 1:1 to the production LLM + CLIP + RAG backend on slide 04.

06 Success Metrics & Projected Business Impact

Every feature is anchored to a specific, measurable KPI tied to a business outcome. Projections are modelled on published benchmarks for AI-driven personalisation, applied to Myntra's FY24 scale.

25%

Return Reduction

FitGenius removes the #1 return driver — fit mismatch — by combining review mining with the user's body data before they buy.

₹525 Cr saved / yr

18%

Conversion Uplift

StyleGenie compresses discovery from 40 items to 5 curated looks, dramatically shortening the path from intent to purchase.

Direct GMV lift

22%

AOV Increase

Outfit Studio shifts buying from a single item to the complete look — an average of +1.8 items per order.

+₹340 / order

2.4x

Engagement

TrendPulse builds a daily check-in habit — users return for AI-curated trends, not just to shop.

DAU / MAU ratio

Projected Impact: Before vs After StyleGenie

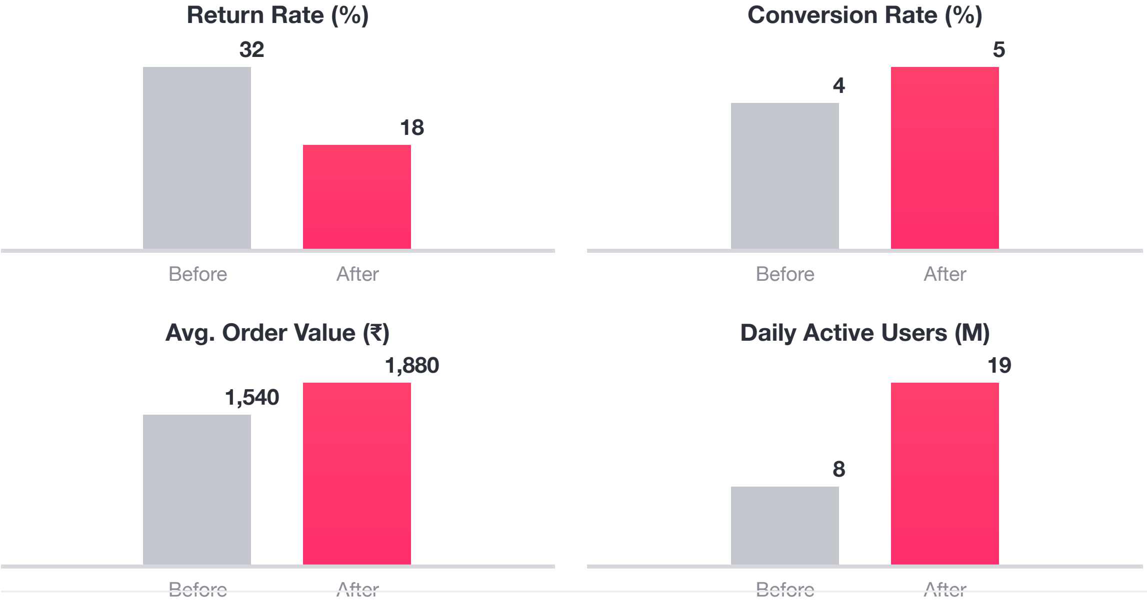

Feature-Level KPI Table

Feature	Metric	Target
StyleGenie Chat	Chat → cart conversion	12%
Visual Match	Search → PDP CTR	35%
Outfit Studio	Items per order	+1.8
FitGenius	Return rate	<18%
TrendPulse	DAU lift	+40%

Total Projected Annual Impact: ₹1,200–1,700 Cr

Revenue uplift ₹800–1,200 Cr + Return savings ₹400–525 Cr + Premium subscription ₹50–80 Cr.

Sources: McKinsey, "The State of Fashion 2024" (return-cost benchmarks); BCG, "AI in Retail" (conversion ranges); Statista, "Myntra Revenue & User Metrics FY24".

07 Risks, Pitfalls & Mitigation Strategy

Every AI product has failure modes — the discipline is to design for them upfront, not discover them in production. These are the six most critical risks for StyleGenie, each rated by severity with a specific, implementable mitigation.

⚠️ Hallucinations — Critical

RISK
AI suggests products that don't exist, are out of stock, or sit in the wrong category — which would destroy user trust instantly.

MITIGATION
Every response is grounded through a **RAG pipeline** — the LLM never invents product data, it only references the live catalog. Confidence scoring filters low-certainty output, and every card links to a real PDP.

🛡️ Privacy — High

RISK
Users are uncomfortable sharing body measurements, photos or style data, and India's DPDP Act imposes strict obligations.

MITIGATION
On-device image processing (CLIP runs locally). Explicit opt-in per data type with a clear benefit. Zero third-party sharing, and full DPDP compliance with deletion on request.

⚡ Latency — High

RISK
LLM inference can take 3–5 seconds — far too slow for a shopping flow where users expect near-instant results.

MITIGATION
Streaming responses (word-by-word) cut perceived wait; pre-computed outfit caches for common queries; edge-deployed CLIP for visual search. Target SLA: under 2 seconds for every AI feature.

🤝 User Trust — Medium

RISK
"How can a machine know what looks good on me?" — adoption is gated on whether users trust the AI's judgement.

MITIGATION
Transparent confidence scores ("94% match"), evidence-based explanations ("based on 10,247 reviews"), and a manual override at every step. The AI earns credibility through accuracy over time.

📈 Scalability — High

RISK
GPT-4o costs at Myntra's 50M+ MAU scale could reach ₹50–100 Cr/month — unsustainable without optimisation.

MITIGATION
Tiered models (Gemini Flash / Haiku for simple queries, premium only for complex styling), aggressive caching so ~100 similar queries share one response, and batch processing for TrendPulse.

👥 Adoption — Medium

RISK
Users default to familiar browse-filter-scroll. AI features get ignored if presented as a separate, unfamiliar experience.

MITIGATION
Soft inline integration — AI appears inside the existing flow, not a separate tab. Contextual nudges fire on struggle signals (long browse, repeated searches). A free tier builds the habit; premium upsell follows.

Sources: "On the Dangers of Stochastic Parrots" (Bender et al., 2021); India DPDP Act 2023; OpenAI Pricing & Rate-Limit Docs; Nielsen Norman Group, "AI Feature Adoption Patterns".