Reimagining Myntra with Generative AI

Myntra StyleGenie — an AI-native styling intelligence layer that transforms fashion discovery from passive browsing into personalised, intent-driven experiences.

Akshat Singh
Roll No: 23113019

01 The Problem: Fashion E-Commerce is Broken

Myntra serves **50M+ monthly active users** across India, offering **9 lakh+ products** from 5,000+ brands. Despite this scale, the platform suffers from fundamental UX inefficiencies that cost the business **over ₹2,000 Cr annually in returns alone**. The root cause: users navigate massive catalogs through keyword search and static filters — tools built for the desktop era, not for how people actually think about fashion.

70%
Cart Abandonment
Users add items but can't decide. Choice overload creates decision paralysis — the average user browses 40+ products before a single purchase.

35%
Return Rate
Fashion has the highest return rate in e-commerce. Primary driver: fit & style mismatch — shoppers can't try on, and size charts are unreliable.

3.2 min
Time to Find
Average time to find a relevant product. Users scroll pages of irrelevant results because search is keyword-based, not intent-aware.

₹2,100 Cr
Annual Return Cost
Reverse logistics, restocking, customer service & damaged goods. Pure cost — it generates zero revenue for the platform.

Who is affected? — User Segments

The Browser (62% of DAU)
Scrolls endlessly, rarely converts. Needs curated discovery — not more filters.

The Returner (31% return rate)
Buys 5, keeps 1. Fit & style mismatch is the core pain. Needs fit intelligence.

The Occasion Shopper (High AOV)
Needs a complete outfit for a wedding/party/fest. No idea where to start. Needs intent curation.

The Trend Chaser (18–28)
Wants celebrity/influencer looks at their budget. Needs visual search & trend matching.

Hypothesis: *A Generative AI styling layer can reduce returns by 25%, lift conversion by 18%, and raise AOV by 22% — creating* **₹1,200–1,700 Cr in annual impact**.

Sources: Myntra FY24 Annual Report (GMV & MAU); RedSeer, "Fashion E-commerce Returns in India, 2024"; Baymard Institute, "Cart Abandonment Statistics, 2024".

02 Why Generative AI? Traditional Approaches Have Hit a Ceiling

Myntra already uses predictive ML (collaborative filtering, content-based recos). These work for simple cases but **fundamentally cannot** handle the problems above. Below: why each traditional approach fails, and how GenAI solves it — these are categorically different capabilities, not incremental improvements.

User Problem	Traditional ML	Why It Fails	GenAI Solution	How It Works
"Outfit for a college fest under ₹2K"	Keyword search + price filter	Can't understand occasion, vibe or styling context — only matches words	**LLM conversational search**	LLM parses intent (occasion + budget + style) → vector retrieval → composes outfits
"Find me something like what Alia wore"	Text-based image tags	Manual tagging can't capture aesthetic — "boho vibes" has no tag	**CLIP multimodal visual search**	CLIP encodes image + text in one space, finds visually similar catalog items
"Will this fit me?"	Static size-chart lookup	Ignores brand variation, fabric stretch, body diversity, user history	**NLP review mining + body profile**	Mines 10K+ reviews for fit signals, cross-refs the user's body data
"Style this kurta into a complete look"	"Frequently bought together"	Co-purchase ≠ styling. No colour, occasion or aesthetic sense	**Generative outfit composition**	LLM applies styling rules, generates occasion-specific outfits from real products
"What should I wear this week?"	"Trending" product list	Same list for everyone. No personalisation to individual style	**AI trend digest + style graph**	LLM generates trend narratives, curates to the user's evolving style profile

Key Insight — Traditional ML predicts what you might *click* next. GenAI understands what you actually *need*. It's the difference between a search engine and a personal stylist — categorically new capabilities enabled by LLMs, multimodal models and generative composition.

Sources: OpenAI CLIP (Radford et al., 2021); "Retrieval-Augmented Generation for Knowledge-Intensive NLP" (Lewis et al., 2020); Myntra Engineering Blog, "Search & Discovery Architecture".

03 My Approach & the StyleGenie Solution

My approach: from problem to solution

01 DIAGNOSE
Studied 4 root pains: choice overload (40+ browses), fit uncertainty (35% returns), context blindness, manual discovery loop.

02 REFRAME
Shifted from inventory-first ("browse 9 lakh products") to intent-first ("look great at a wedding under ₹5K").

03 DESIGN
Mapped each segment + pain to one specific GenAI capability. Each feature solves exactly one root pain — no overlap.

04 VALIDATE
Anchored every feature to a measurable KPI (return rate, conversion, AOV, DAU). No feature ships without a metric.

The solution: 5 GenAI features (each mapped to a user pain)

💬 StyleGenie Chat
AI Concierge
Solves: Choice overload

WHAT
Conversational AI that understands occasion, budget, body & style via natural language.

HOW
LLM + RAG: intent parse → vector search → outfit compose → stream.

VALUE
40 products → 5 outfits. 12% chat-to-cart.

📷 Visual Match
Snap & Find
Solves: Style discovery

WHAT
Upload any photo (screenshot, Pinterest) → AI finds visually similar Myntra products.

HOW
CLIP encodes image → cosine similarity vs catalog embeddings → ranked results.

VALUE
35% CTR (vs 8% keyword search).

🪄 Outfit Studio
Composer
Solves: No outfit thinking

WHAT
Pick one product → AI generates 3–5 complete styled outfits using real catalog items.

HOW
Style-attribute encoding → LLM applies colour theory + occasion rules + composition.

VALUE
+1.8 items/order. 22% AOV lift.

🎯 FitGenius
Fit Predictor
Solves: Fit uncertainty

WHAT
Predicts your perfect size with 90%+ confidence using review mining + body profile.

HOW
NLP extracts fit signals from 10K reviews → cross-ref with user's body data.

VALUE
Returns 32% → <18%. ~₹525 Cr saved/yr.

📈 TrendPulse
Trend Digest
Solves: Static discovery

WHAT
Personalised daily fashion feed with AI-curated trend lookbooks for your style.

HOW
Social-signal monitor → LLM trend narratives → products matched to trend + style graph.

VALUE
+40% DAU from daily engagement.

User segment → primary features: Browser → Chat + TrendPulse · Returner → FitGenius + Studio · Occasion → Chat + Studio · Trend Chaser → Visual + Pulse

Sources: CLIP — Radford et al. (OpenAI, 2021); RAG — Lewis et al. (Facebook AI, 2020); Myntra Category Report FY24.

04 Implementation Flowchart: How StyleGenie Actually Works

End-to-end system flow: user query → AI orchestration → model layer → data retrieval → personalised response. Colour-coded by system layer.

① USER

"Outfit for Diwali under ₹3K, traditional but modern"

▼

② API GATEWAY + LLM ROUTER

Auth → Rate-limit → Intent Classifier (LLM) extracts {occasion, budget, style, body} and dispatches to the right pipeline(s)

▼ ▼ ▼ ▼ ▼

③ AI PIPELINES

StyleGenie Chat	Visual Match	Outfit Studio	FitGenius	TrendPulse
GPT-4o + RAG	CLIP-ViT	GPT-4o + Style Engine	NLP Review Miner	LLM + Trend Feed
Intent parse → vector search → outfit compose → stream	Encode image → cosine similarity → top-k catalog match	Item attrs → colour theory → occasion rules → 5 outfits	Mine reviews → fit signals → cross-ref body → size prediction	Social signal → narrative gen → style-graph match

▼

④ SHARED DATA LAYER (RAG grounding — the LLM only ever references live data)

Pinecone Vector DB	Product Catalog	Review Corpus	User Style Graph
9 lakh+ product embeddings	SKUs, prices, live inventory	10M+ reviews, fit-indexed	per-user preference vectors

▼

⑤ **RESPONSE SYNTHESIS** Compose outfit cards • add confidence scores • cite review evidence • stream

⑥ **USER SEES OUTPUT** 3–5 Diwali outfit cards with prices, fit confidence & "Add to Bag"

Sources: "Attention Is All You Need" (Vaswani et al., 2017); Pinecone Vector DB Architecture; OpenAI Function-Calling Docs; Google Gemini API Reference; Anthropic RAG Pattern Guide.

05 The Live Prototype: How Each Feature Works In the App

A working React prototype rebuilt in Myntra's UI, with all five GenAI features interactive. Below is the **real in-app interaction** for each — what the user taps and what appears on screen — beside screenshots from the running app.

💬 **StyleGenie Chat —** tap the floating "✨ **Ask StyleGenie**" button → type *"festive look under 3k for women"* → request appears as **intent chips** (festive · ≤₹3,000 · women) → **5 outfit cards stream in** → tap a card to open its product page.

📷 **Visual Match —** tap the 📷 **in the search bar** → upload an image / pick a vibe → an **"encoding with CLIP…"** animation → a grid of products with **% match scores**.

🪄 **Outfit Studio —** on any product page tap **"Generate complete looks"** → **3–5 styled looks** (top + bottom + footwear + accessory) each with a price total & occasion → **"Add complete look to bag"**.

🎯 **FitGenius —** enter height / weight / usual size / build → tap **"Find my size"** → size shown with a **confidence bar + review evidence** ("runs slim — we sized up").

📈 **TrendPulse —** open the **TrendPulse** tab → scroll **AI-written trend cards** (e.g. "Quiet Luxury"), each with a curated product row.

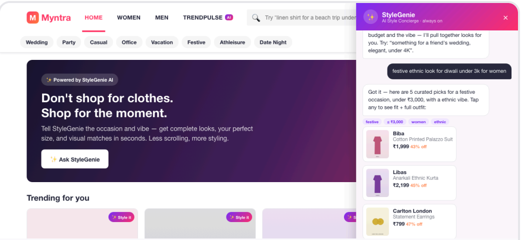

StyleGenie Chat — intent chips + streamed cards

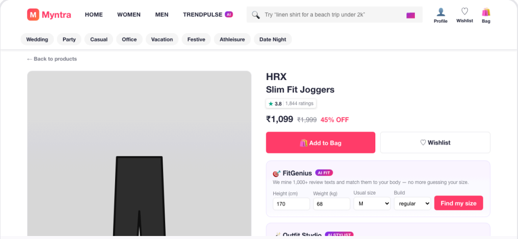

Product page — FitGenius + Outfit Studio inline

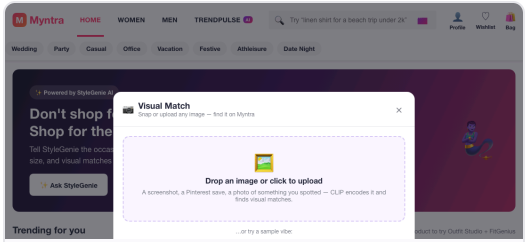

Visual Match — CLIP-style image search

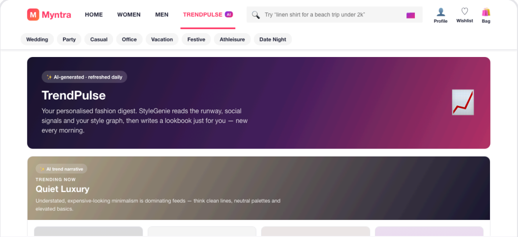

TrendPulse — AI trend narratives

Interaction model: a floating concierge + inline AI cards — discoverable, never blocking the core flow. **Try it live:** deep links `#genie/festive%20look%20under%203k`, `#p/12`, `#visual`, `#trends` jump to each feature. **Stack:** React + Vite; AI runs client-side for the demo, mapping 1:1 to the production LLM + CLIP + RAG backend in slide 04.

06 Success Metrics & Projected Business Impact

Each feature maps to a specific, measurable KPI tied to a business outcome. Projections use industry benchmarks for AI-driven personalisation in e-commerce.

25%
Return Reduction
FitGenius eliminates the #1 return driver (fit mismatch) via review mining + body data.

₹525 Cr saved/yr

18%
Conversion Uplift
StyleGenie cuts discovery from 40 items to 5 curated looks — shortening path to purchase.

Direct GMV lift

22%
AOV Increase
Outfit Studio drives bundle buying: the complete look instead of one item (+1.8 items/order).

+₹340/order

2.4x
Engagement
TrendPulse creates daily check-in behaviour — users return for AI-curated trends, not just to shop.

DAU/MAU ratio

Projected Impact: Before vs After

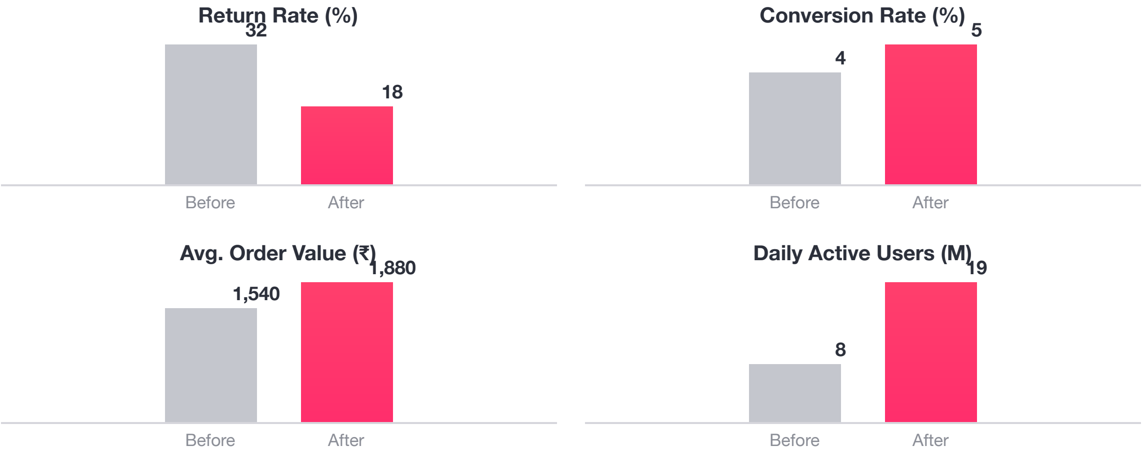

Feature-Level KPI Table

Feature	Metric	Target
StyleGenie Chat	Chat → cart	12%
Visual Match	Search CTR	35%
Outfit Studio	Items / order	+1.8
FitGenius	Return rate	<18%
TrendPulse	DAU lift	+40%

Total Projected Annual Impact: ₹1,200–1,700 Cr
Revenue uplift ₹800–1,200 Cr + Return savings ₹400–525 Cr + Premium subscription ₹50–80 Cr.

Sources: McKinsey, "The State of Fashion 2024" (return-cost benchmarks); BCG, "AI in Retail" (conversion ranges); Statista, "Myntra Revenue & User Metrics FY24".

07 Risks, Pitfalls & Mitigation Strategy

Every AI product has failure modes. The key is not avoiding them — it's designing for them upfront. Below are the six most critical risks for StyleGenie, each with a specific, implementable mitigation.

⚠️ Hallucinations

Critical

RISK
AI suggests products that don't exist, are out of stock, or in the wrong category — destroying trust instantly.

MITIGATION
All responses grounded via a **RAG pipeline** — the LLM never invents product data, only references the live catalog. Confidence scoring filters low-certainty output; every card links to a real PDP.

🛡️ Privacy

High

RISK
Users uncomfortable sharing body measurements, photos or style data. GDPR / India DPDP Act compliance.

MITIGATION
On-device image processing (CLIP runs locally). Explicit opt-in per data type. Zero third-party sharing. Full DPDP compliance with deletion on request.

⚡ Latency

High

RISK
LLM inference takes 3–5s — too slow for a shopping context where users expect instant results.

MITIGATION
Streaming responses (word-by-word reduces perceived wait). Pre-computed outfit caches. Edge-deployed CLIP. Target SLA <2s for all AI features.

🤝 User Trust

Medium

RISK
"How can a machine know what looks good on me?" Adoption depends on trust.

MITIGATION
Transparent confidence scores ("94% match"). Evidence-based explanations ("based on 10,247 reviews"). Manual override at every step. Trust earned through accuracy over time.

📈 Scalability

High

RISK
GPT-4o costs at 50M+ MAU could reach ₹50–100 Cr/month — unsustainable without optimisation.

MITIGATION
Tiered models (Gemini Flash/Haiku for simple, premium only for complex styling). Aggressive caching — 100 similar queries share one response. Batch processing for TrendPulse.

👥 Adoption

Medium

RISK
Users stick to browse-filter-scroll. AI features get ignored if presented as a separate, unfamiliar experience.

MITIGATION
Soft inline integration — AI appears within the existing flow, not a separate tab. Contextual nudges on struggle signals (long browse, repeated searches). Free tier to build habit, premium upsell later.

Sources: "On the Dangers of Stochastic Parrots" (Bender et al., 2021); India DPDP Act 2023; OpenAI Pricing & Rate-Limit Docs; Nielsen Norman Group, "AI Feature Adoption Patterns".